

part2

December 5, 2022

1 Part 2: Illness Detection

1.0.1 This part determines level of illness within sheep based on sitting/standing ground truth data

```
[ ]: # Imports
import pandas as pd
import numpy as np
import statistics
from sklearn.preprocessing import StandardScaler

[ ]: # read in the data first
# note that D is missing several days
files = ['A.csv', 'B.csv', 'U.csv', 'V.csv', 'W.csv', 'D.csv']
sheep_raw = []
data_dir = 'data/'
for file in files:
    file = data_dir + file
    df = pd.read_csv(file)
    # drop unnecessary columns
    df = df.drop(columns=df.columns[3:])
    df = df.drop(columns=df.columns[0:2])

    # rename column
    df.columns = ['Standing']

    # change accelerometer value to 1 if standing, 0 if sitting
    # THERE'S A SPECIAL CASE FOR SHEEP V and F, THE ACCELEROMETER WAS UPSIDE_
    ↪DOWN
    if file == 'data/V.csv' or file == 'data/F.csv':
        df['Standing'] = np.where(df['Standing'] > 0.49, 1, 0)
    else:
        df['Standing'] = np.where(df['Standing'] < -0.4, 1, 0)
    sheep_raw.append(df)

[ ]: # version of reading in ALL data
```

```

all_files = ['A.csv', 'B.csv', 'C.csv', 'D.csv', 'E.csv', 'F.csv', 'G.csv', 'H.
↪csv', 'I.csv', 'J.csv', 'K.csv', 'L.csv', 'M.csv', 'N.csv', 'O_0.csv', 'P.
↪csv', 'Q.csv', 'R_0.csv', 'S.csv', 'T.csv', 'U.csv', 'V.csv', 'W.csv']

sheep_raw_full = []
data_dir = 'data/'
for file in all_files:
    file = data_dir + file
    df = pd.read_csv(file)
    # drop unnecessary columns
    df = df.drop(columns=df.columns[3:])
    df = df.drop(columns=df.columns[0:2])

    # rename column
    df.columns = ['Standing']

    # change accelerometer value to 1 if standing, 0 if sitting
    # THERE'S A SPECIAL CASE FOR SHEEP V and F, THE ACCELEROMETER WAS UPSIDE_
↪DOWN
    if file == 'data/V.csv' or file == 'data/F.csv':
        df['Standing'] = np.where(df['Standing'] > 0.49, 1, 0)
    else:
        df['Standing'] = np.where(df['Standing'] < -0.4, 1, 0)
    sheep_raw_full.append(df)

```

```

[ ]: for sheep, file in zip(sheep_raw_full, all_files):
    temp = sheep['Standing'].to_list()
    total_sit = 0
    total_stand = 0
    for instance in temp:
        total_stand += instance
        if instance == 0:
            total_sit += 1
    if total_sit/(total_sit+total_stand) > 0.9:
        print(f'{file}: {total_stand}')
    print(f'Sit percentage: {round(total_sit/(total_sit+total_stand), 4) *
↪100}%\nTotal Size: {total_stand+total_sit}\n')

```

Sit percentage: 63.01%
Total Size: 21686

Sit percentage: 62.67%
Total Size: 21686

Sit percentage: 63.6%
Total Size: 21686

Sit percentage: 57.4%
Total Size: 5831

Sit percentage: 64.57000000000001%
Total Size: 21686

Sit percentage: 79.56%
Total Size: 21686

Sit percentage: 61.82%
Total Size: 21686

Sit percentage: 65.99000000000001%
Total Size: 21686

Sit percentage: 52.92%
Total Size: 21686

Sit percentage: 56.64%
Total Size: 21686

Sit percentage: 49.27%
Total Size: 21686

Sit percentage: 70.37%
Total Size: 21686

Sit percentage: 72.84%
Total Size: 21686

Sit percentage: 63.93%
Total Size: 21686

Sit percentage: 63.84999999999994%
Total Size: 21686

Sit percentage: 72.22%
Total Size: 2977

Sit percentage: 53.38%
Total Size: 21686

Sit percentage: 65.38000000000001%
Total Size: 21686

Sit percentage: 49.25%
Total Size: 21686

Sit percentage: 69.179999999999999%
Total Size: 21686

Sit percentage: 64.3%
Total Size: 21686

Sit percentage: 64.92%
Total Size: 21686

Sit percentage: 63.7%
Total Size: 21686

```
[ ]: print(sheep_raw[0].shape)
      print(sheep_raw[1].shape)
      print(sheep_raw[2].shape)
      print(sheep_raw[3].shape)
      print(sheep_raw[4].shape)
      print(sheep_raw[5].shape)
```

```
(21686, 1)
(21686, 1)
(21686, 1)
(21686, 1)
(21686, 1)
(5831, 1)
```

1.0.2 Windowing

```
[ ]: # Function to window our samples
      def splitSamples(df: pd.DataFrame, samples_per_split: int) -> list:
          all_samples = []
          start_index = 0
          end_index = samples_per_split
          while end_index < len(df):
              temp_df = df[start_index:end_index]
              all_samples.append(temp_df)
              start_index += samples_per_split // 2
              end_index += samples_per_split // 2

          return all_samples
```

```
[ ]: # Call the function with each dataframe to get sample lists

      # The accelerometer records 2880 times a day, so let's try a window of that size
      SAMPLES_PER_SPLIT = 2880
      for i in range(len(sheep_raw)):
```

```

sheep_raw[i] = splitSamples(sheep_raw[i], SAMPLES_PER_SPLIT)

for i in range(len(sheep_raw_full)):
    sheep_raw_full[i] = splitSamples(sheep_raw_full[i], SAMPLES_PER_SPLIT)

```

```

[ ]: # Convert dataframes into just a list of sit/stand
for i in range(len(sheep_raw)):
    for j in range(len(sheep_raw[i])):
        sheep_raw[i][j] = sheep_raw[i][j]['Standing'].tolist()

# Convert dataframes into just a list of sit/stand
for i in range(len(sheep_raw_full)):
    for j in range(len(sheep_raw_full[i])):
        sheep_raw_full[i][j] = sheep_raw_full[i][j]['Standing'].tolist()

```

1.0.3 Feature Engineering

```

[ ]: def extractFeatures(windows: list) -> list:
    """
    There are a few features we should extract here
    1. total sits
    2. total stands
    3. changes (how many times did it go from sit to stand or vice versa)
    4. mean (essentially what percentage of the time was it standing)
    5. mode (was it standing or sitting more)
    6. longest sit
    7. longest stand
    8. Average length of sitting bouts
    9. Average length of standing bouts
    10. Std dev of sitting bouts
    11. Standard deviation of standing bouts
    """
    X = []

    for window in windows:
        temp = []
        total_sit = 0
        total_stand = 0
        total_changes = 0
        longest_sit = 0
        longest_stand = 0
        average_sit = []
        average_stand = []

        current_sit_streak = 0
        current_stand_streak = 0

```

```

# mean and median can be calculated without looping
mean = np.mean(window)
mode = statistics.mode(window)

prev_state = window[0]

for sample in window:
    if sample == 0:
        total_sit += 1
        current_sit_streak += 1
        if current_sit_streak > longest_sit:
            longest_sit = current_sit_streak
        if prev_state != sample:
            total_changes += 1
            average_stand.append(current_stand_streak)
            current_stand_streak = 0
    else:
        total_stand += 1
        current_stand_streak += 1
        if current_stand_streak > longest_stand:
            longest_stand = current_stand_streak
        if prev_state != sample:
            total_changes += 1
            average_sit.append(current_sit_streak)
            current_sit_streak = 0

    # add on any leftover sits or stands
    if current_sit_streak != 0:
        average_sit.append(current_sit_streak)
    else:
        average_stand.append(current_stand_streak)

    # can't have empty lists
    if not average_sit:
        average_sit.append(0)
    if not average_stand:
        average_stand.append(0)
    prev_state = sample

temp.append(total_sit)
temp.append(total_stand)
temp.append(total_changes)
temp.append(mean)
temp.append(mode)
temp.append(longest_sit)
temp.append(longest_stand)
temp.append(np.mean(average_sit))

```

```

        temp.append(np.mean(average_stand))
        temp.append(np.std(average_sit))
        temp.append(np.std(average_stand))

    X.append(temp.copy())

return X

```

```

[ ]: sheep = []
    for i in range(len(sheep_raw)):
        sheep.append(extractFeatures(sheep_raw[i]))

    sheep_full = []
    for i in range(len(sheep_raw_full)):
        sheep_full.append(extractFeatures(sheep_raw_full[i]))

```

```

[ ]: print(sheep[0][0])
    print(sheep[0][1])

    print(sheep_full[0][0])
    print(sheep_full[0][1])

```

```

[1999, 881, 118, 0.3059027777777778, 0, 175, 116, 62.21137026239067,
22.103081827842722, 42.75723646511831, 24.541816253464393]
[1885, 995, 114, 0.3454861111111111, 0, 198, 116, 58.02368692070031,
25.64577397910731, 43.31277351234201, 26.867492965424038]
[1999, 881, 118, 0.3059027777777778, 0, 175, 116, 62.21137026239067,
22.103081827842722, 42.75723646511831, 24.541816253464393]
[1885, 995, 114, 0.3454861111111111, 0, 198, 116, 58.02368692070031,
25.64577397910731, 43.31277351234201, 26.867492965424038]

```

1.0.4 Create X and Y to prepare model

```

[ ]: X = []
    X_RAW = []
    y = []

    # Sheep have 5 levels of illness. We group these into 3 levels:
    # 1 = not sick
    # 2,3 = moderately sick
    # 4,5 = very sick
    # Sheep A, B, U, V, W, and D have sickness levels of 4, 3, 5, 1, 2, and 5
    # respectively
    SHEEP_ILLNESS_LEVEL = [2, 1, 2, 0, 1, 2]
    for s, sr, ill in zip(sheep, sheep_raw, SHEEP_ILLNESS_LEVEL):
        for window, window_raw in zip(s, sr):
            X.append(window)

```

```

        X_RAW.append(window_raw)
        y.append(ill)

X_FULL = []
X_RAW_FULL = []
y_full = []
# FULL_SHEEP_ILLNESS_LEVEL = [4, 3, 1, 5, 2, 1, 3, 5, 4, 1, 4, 4, 2, 2, 3, 3,
↪5, 3, 4, 1, 5, 1, 2]
FULL_SHEEP_ILLNESS_LEVEL = [2, 1, 0, 2, 1, 0, 1, 2, 2, 0, 2, 2, 1, 1, 1, 1, 2,
↪1, 2, 0, 2, 0, 1]
for s, sr, ill in zip(sheep_full, sheep_raw_full, FULL_SHEEP_ILLNESS_LEVEL):
    for window, window_raw in zip(s, sr):
        X_FULL.append(window)
        X_RAW_FULL.append(window_raw)
        y_full.append(ill)

y_full_binary = []
# FULL_SHEEP_BINARY_ILLNESS_LEVEL = [4, 3, 1, 5, 2, 1, 3, 5, 4, 1, 4, 4, 2, 2,
↪3, 3, 5, 3, 4, 1, 5, 1, 2]
FULL_SHEEP_BINARY_ILLNESS_LEVEL = [1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
↪1, 1, 1, 1, 0, 1, 0, 0]
for s, ill in zip(sheep_full, FULL_SHEEP_BINARY_ILLNESS_LEVEL):
    for window in s:
        y_full_binary.append(ill)

# normalize X
sc = StandardScaler()
X = sc.fit_transform(X)
sc = StandardScaler()
X_FULL = sc.fit_transform(X_FULL)

```

1.0.5 Model Creation

```

[ ]: # Train and evaluate the model
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import numpy as np

def train_and_evaluate_model(X, y, model, binary):
    # Do the K-fold cross validation
    cv = StratifiedKFold(n_splits=10)
    scores = cross_val_score(model, X, y, cv=cv, scoring='accuracy')

    # This function sets an empty array of size len(y)

```



```

    # It then does the predictions for each fold (retraining the model each
    ↪time)
    # and fills in the prediction in the correct spot in the array. By the end
    ↪of it,
    # you will have your y array and a new array (predictions) of guesses for
    ↪each sample
    # this is because Kfold uses each sample once for test data.
    predictions = cross_val_predict(model, X, y, cv=cv)

    # we can use this to create a confusion matrix
    if binary == 0:
        new_labels = ['Not Sick', 'Moderately Sick', 'Very Sick']
        old_labels = [0, 1, 2]
    else:
        new_labels = ['Healthy', 'Sick']
        old_labels = [0, 1]

    cm = confusion_matrix(y, predictions)
    plt.matshow(cm, cmap=plt.cm.Blues)
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            plt.text(x=j, y=i, s=cm[i, j], va='center', ha='center',
            ↪size='xx-large')
    plt.xticks(ticks=old_labels, labels=new_labels)
    plt.yticks(ticks=old_labels, labels=new_labels)
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Label')

    print(f'All accuracy scores from 10 folds:')
    for score in scores:
        print(f'{round(score * 100, 2)}%', end=' ')
    print()
    print(f'Mean Accuracy from 10 Fold Cross-Validation: {round(np.
    ↪mean(scores)*100, 2)}%\n')

```

1.0.6 Testing Models

```

[ ]: from sklearn.ensemble import RandomForestClassifier

print('USING RAW DATA')
model = RandomForestClassifier(n_estimators=100, random_state=42)
train_and_evaluate_model(X_RAW, y, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA')
model = RandomForestClassifier(n_estimators=100, random_state=42)

```

```

train_and_evaluate_model(X, y, model, 0)

print('*****')
print('USING RAW DATA WITH ALL SHEEP')
model = RandomForestClassifier(n_estimators=100, random_state=42)
train_and_evaluate_model(X_RAW_FULL, y_full, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP')
model = RandomForestClassifier(n_estimators=100, random_state=42)
train_and_evaluate_model(X_FULL, y_full, model, 0)

print('*****')
print('USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = RandomForestClassifier(n_estimators=100, random_state=42)
train_and_evaluate_model(X_RAW_FULL, y_full_binary, model, 1)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = RandomForestClassifier(n_estimators=100, random_state=42)
train_and_evaluate_model(X_FULL, y_full_binary, model, 1)

```

USING RAW DATA

All accuracy scores from 10 folds:

50.0% 37.5% 50.0% 71.43% 42.86% 42.86% 42.86% 57.14% 42.86% 0.0%

Mean Accuracy from 10 Fold Cross-Validation: 43.75%

USING FEATURE ENGINEERED DATA

All accuracy scores from 10 folds:

37.5% 12.5% 37.5% 71.43% 85.71% 71.43% 71.43% 71.43% 100.0% 71.43%

Mean Accuracy from 10 Fold Cross-Validation: 63.04%

USING RAW DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

40.0% 33.33% 33.33% 53.33% 36.67% 30.0% 46.67% 50.0% 41.38% 34.48%

Mean Accuracy from 10 Fold Cross-Validation: 39.92%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

36.67% 46.67% 36.67% 36.67% 43.33% 60.0% 53.33% 40.0% 44.83% 37.93%

Mean Accuracy from 10 Fold Cross-Validation: 43.61%

USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT

All accuracy scores from 10 folds:

60.0% 66.67% 50.0% 56.67% 63.33% 53.33% 73.33% 43.33% 51.72% 62.07%

Mean Accuracy from 10 Fold Cross-Validation: 58.05%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT


All accuracy scores from 10 folds:


70.0% 76.67% 76.67% 60.0% 40.0% 53.33% 66.67% 76.67% 48.28% 75.86%

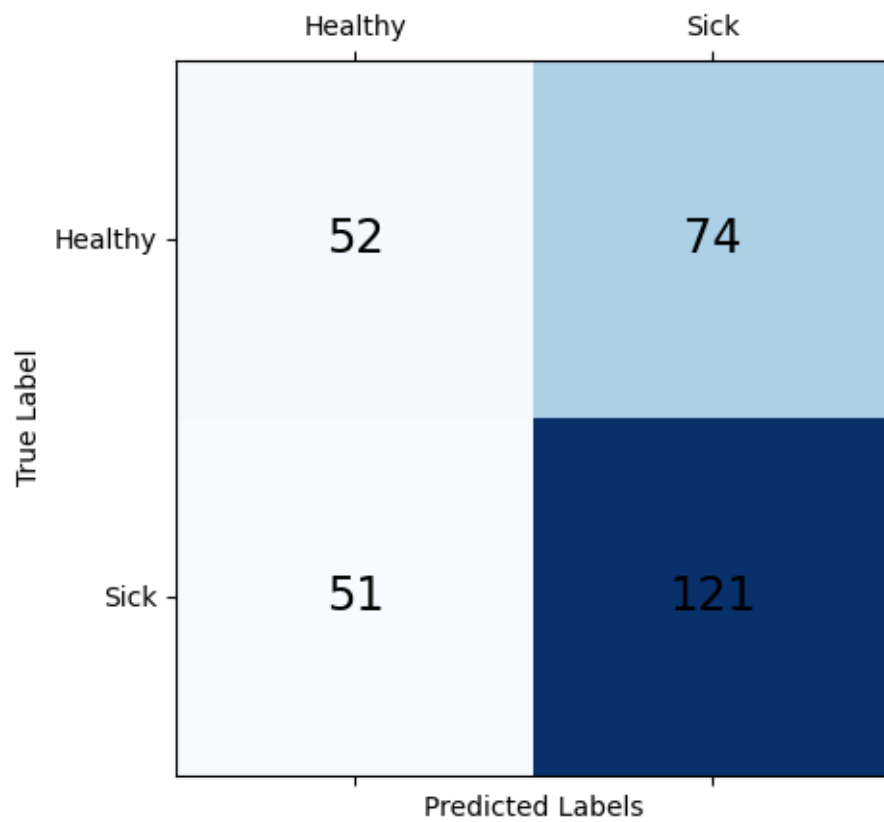
Mean Accuracy from 10 Fold Cross-Validation: 64.41%

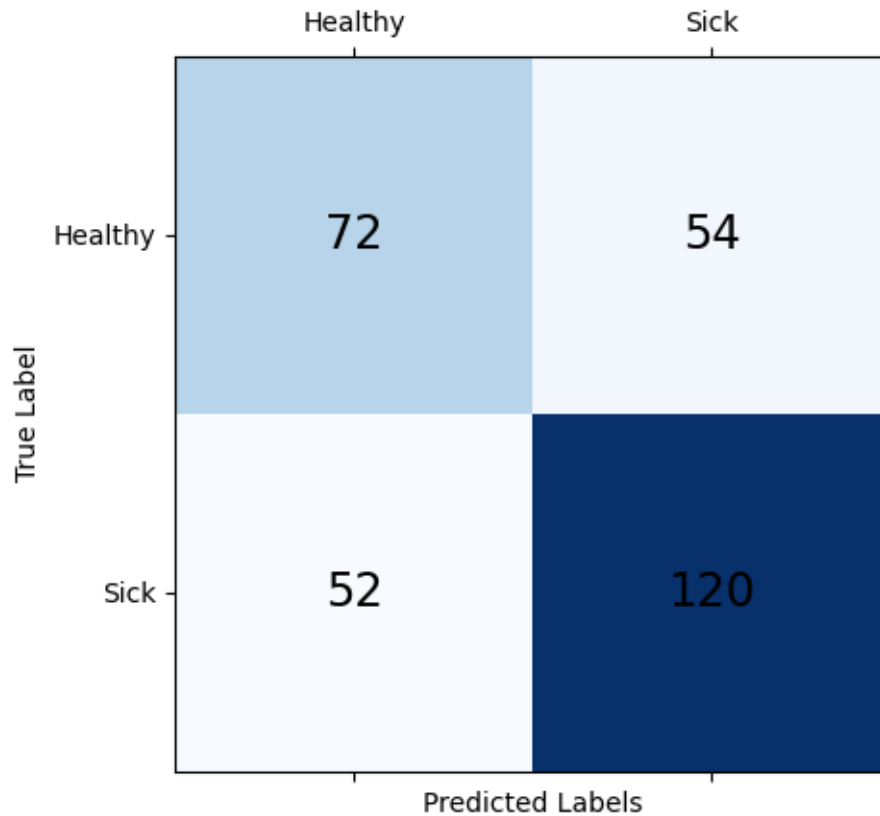
		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	2	3	9
	Moderately Sick	1	10	17
	Very Sick	2	9	20
		Predicted Labels		

		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	8	5	1
	Moderately Sick	3	16	9
	Very Sick	3	7	21
		Predicted Labels		









```
[ ]: from xgboost import XGBClassifier

print('USING RAW DATA')
model = XGBClassifier()
train_and_evaluate_model(X_RAW, y, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA')
model = XGBClassifier()
train_and_evaluate_model(X, y, model, 0)

print('*****')
print('USING RAW DATA WITH ALL SHEEP')
model = XGBClassifier()
train_and_evaluate_model(X_RAW_FULL, y_full, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP')
model = XGBClassifier()
train_and_evaluate_model(X_FULL, y_full, model, 0)
```



```

print('*****')
print('USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = XGBClassifier()
train_and_evaluate_model(X_RAW_FULL, y_full_binary, model, 1)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = XGBClassifier()
train_and_evaluate_model(X_FULL, y_full_binary, model, 1)

```

USING RAW DATA

All accuracy scores from 10 folds:

25.0% 50.0% 50.0% 28.57% 14.29% 42.86% 28.57% 42.86% 42.86% 28.57%

Mean Accuracy from 10 Fold Cross-Validation: 35.36%

USING FEATURE ENGINEERED DATA

All accuracy scores from 10 folds:

50.0% 12.5% 37.5% 28.57% 85.71% 85.71% 85.71% 57.14% 85.71% 71.43%

Mean Accuracy from 10 Fold Cross-Validation: 60.0%

USING RAW DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

50.0% 30.0% 30.0% 30.0% 50.0% 30.0% 50.0% 26.67% 37.93% 24.14%

Mean Accuracy from 10 Fold Cross-Validation: 35.87%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

30.0% 43.33% 40.0% 46.67% 40.0% 70.0% 50.0% 33.33% 51.72% 27.59%

Mean Accuracy from 10 Fold Cross-Validation: 43.26%

USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT

All accuracy scores from 10 folds:

50.0% 53.33% 50.0% 60.0% 63.33% 53.33% 70.0% 60.0% 48.28% 55.17%

Mean Accuracy from 10 Fold Cross-Validation: 56.34%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT


All accuracy scores from 10 folds:


63.33% 60.0% 70.0% 53.33% 43.33% 50.0% 70.0% 70.0% 51.72% 75.86%


Mean Accuracy from 10 Fold Cross-Validation: 60.76%


		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	3	3	8
	Moderately Sick	4	7	17
	Very Sick	4	11	16
		Predicted Labels		

		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	7	5	2
	Moderately Sick	2	17	9
	Very Sick	2	10	19
		Predicted Labels		









```
[ ]: from sklearn.svm import SVC

print('USING RAW DATA')
model = SVC()
train_and_evaluate_model(X_RAW, y, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA')
model = SVC()
train_and_evaluate_model(X, y, model, 0)

print('*****')
print('USING RAW DATA WITH ALL SHEEP')
model = SVC()
train_and_evaluate_model(X_RAW_FULL, y_full, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP')
model = SVC()
train_and_evaluate_model(X_FULL, y_full, model, 0)
```

```

print('*****')
print('USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = SVC()
train_and_evaluate_model(X_RAW_FULL, y_full_binary, model, 1)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = SVC()
train_and_evaluate_model(X_FULL, y_full_binary, model, 1)

```

USING RAW DATA

All accuracy scores from 10 folds:

37.5% 37.5% 37.5% 42.86% 42.86% 42.86% 42.86% 42.86% 42.86% 42.86%

Mean Accuracy from 10 Fold Cross-Validation: 41.25%

USING FEATURE ENGINEERED DATA

All accuracy scores from 10 folds:

50.0% 25.0% 37.5% 85.71% 57.14% 57.14% 57.14% 71.43% 100.0% 71.43%

Mean Accuracy from 10 Fold Cross-Validation: 61.25%

USING RAW DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

46.67% 46.67% 43.33% 56.67% 66.67% 40.0% 56.67% 53.33% 51.72% 44.83%

Mean Accuracy from 10 Fold Cross-Validation: 50.66%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

36.67% 50.0% 36.67% 50.0% 63.33% 43.33% 63.33% 56.67% 65.52% 41.38%

Mean Accuracy from 10 Fold Cross-Validation: 50.69%

USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT

All accuracy scores from 10 folds:

63.33% 76.67% 50.0% 60.0% 63.33% 56.67% 80.0% 60.0% 48.28% 58.62%

Mean Accuracy from 10 Fold Cross-Validation: 61.69%


USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT


All accuracy scores from 10 folds:


60.0% 76.67% 63.33% 50.0% 50.0% 46.67% 80.0% 76.67% 48.28% 79.31%


Mean Accuracy from 10 Fold Cross-Validation: 63.09%


		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	0	0	14
	Moderately Sick	0	0	28
	Very Sick	0	1	30
		Predicted Labels		











```
[ ]: from sklearn.neighbors import KNeighborsClassifier

print('USING RAW DATA')
model = KNeighborsClassifier()
train_and_evaluate_model(X_RAW, y, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA')
model = KNeighborsClassifier()
train_and_evaluate_model(X, y, model, 0)

print('*****')
print('USING RAW DATA WITH ALL SHEEP')
model = KNeighborsClassifier()
train_and_evaluate_model(X_RAW_FULL, y_full, model, 0)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP')
model = KNeighborsClassifier()
train_and_evaluate_model(X_FULL, y_full, model, 0)
```

```

print('*****')
print('USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = KNeighborsClassifier()
train_and_evaluate_model(X_RAW_FULL, y_full_binary, model, 1)

print('')
print('USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT')
model = KNeighborsClassifier()
train_and_evaluate_model(X_FULL, y_full_binary, model, 1)

```

USING RAW DATA

All accuracy scores from 10 folds:

50.0% 37.5% 75.0% 57.14% 14.29% 28.57% 14.29% 28.57% 0.0% 28.57%

Mean Accuracy from 10 Fold Cross-Validation: 33.39%

USING FEATURE ENGINEERED DATA

All accuracy scores from 10 folds:

50.0% 25.0% 50.0% 57.14% 57.14% 57.14% 57.14% 71.43% 100.0% 85.71%

Mean Accuracy from 10 Fold Cross-Validation: 61.07%

USING RAW DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

16.67% 33.33% 40.0% 26.67% 36.67% 33.33% 43.33% 30.0% 41.38% 31.03%

Mean Accuracy from 10 Fold Cross-Validation: 33.24%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP

All accuracy scores from 10 folds:

23.33% 40.0% 43.33% 60.0% 50.0% 30.0% 46.67% 43.33% 41.38% 31.03%

Mean Accuracy from 10 Fold Cross-Validation: 40.91%

USING RAW DATA WITH ALL SHEEP WITH BINARY OUTPUT

All accuracy scores from 10 folds:

50.0% 50.0% 43.33% 46.67% 56.67% 53.33% 43.33% 26.67% 24.14% 65.52%

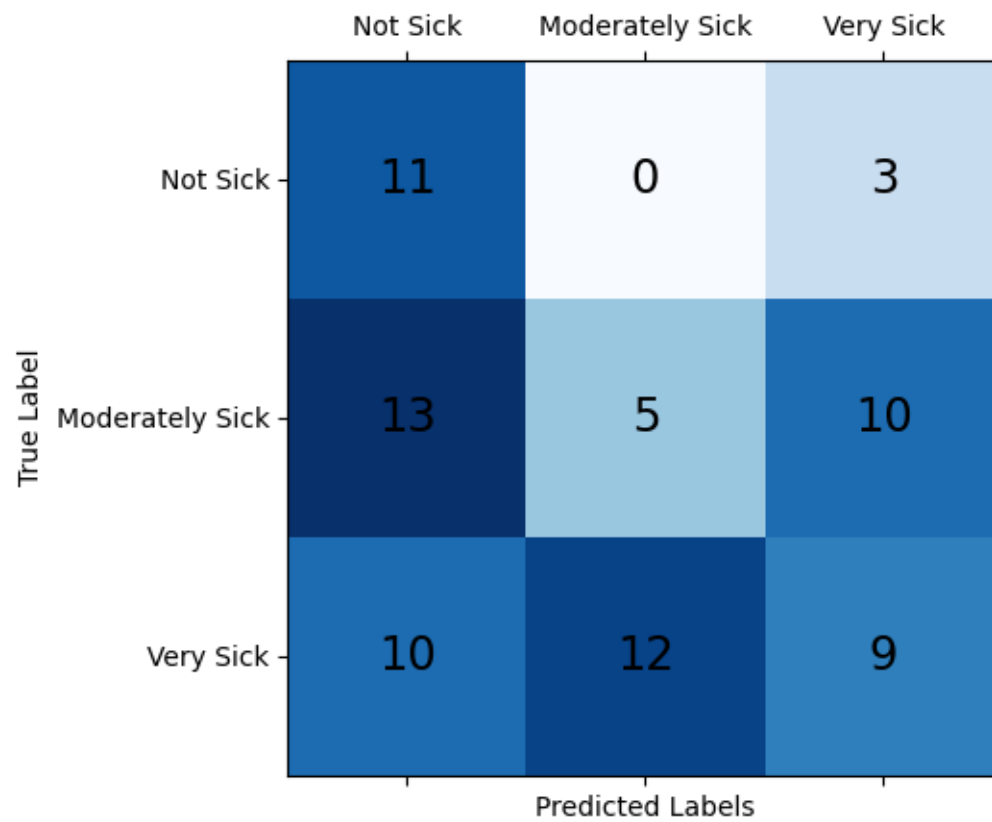
Mean Accuracy from 10 Fold Cross-Validation: 45.97%

USING FEATURE ENGINEERED DATA WITH ALL SHEEP WITH BINARY OUTPUT

All accuracy scores from 10 folds:

70.0% 60.0% 70.0% 50.0% 56.67% 40.0% 60.0% 60.0% 44.83% 65.52%

Mean Accuracy from 10 Fold Cross-Validation: 57.7%



		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	8	2	4
	Moderately Sick	5	20	3
	Very Sick	5	10	16
		Predicted Labels		

		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	23	39	8
	Moderately Sick	51	56	6
	Very Sick	51	44	20
		Predicted Labels		

		Not Sick	Moderately Sick	Very Sick
True Label	Not Sick	20	29	21
	Moderately Sick	47	47	19
	Very Sick	22	38	55
		Predicted Labels		

